

# Malloc Lab

Implementing your own  
malloc/free/realloc

---

# libc malloc/free

---

`void *malloc(size_t size)`

- Allocate *size* bytes and return a pointer to the address allocated address
- `my_type *my_obj = (my_type *)malloc(sizeof(my_type));`

`void free(void *ptr)`

- Free the memory space pointed by *ptr*
- `free(my_obj);`

For more detail,

<https://linux.die.net/man/3/malloc>

# What is malloc?

---

- malloc is designed to provide a simple and portable way to allocate/deallocate a memory block of desired size
- Linux kernel itself also provides very limited dynamic memory management primitives (brk, sbrk)
  - They can only expand/shrink the end of data segment (just like a stack)
- libc, a user-level library, provides malloc implementation using those primitives

# Challenges in malloc design

---

## Execution speed

- Finding a free memory block
- Releasing a memory block

## Memory space consumption

- Data structure overhead
- Internal fragmentation
- External fragmentation

→ Therefore, many different algorithms are there

# Speed evaluation

---

Remember? Big-O notation

- Bubble sort:  $O(n^2)$
- Merge sort:  $O(n \log n)$
- Linear search:  $O(n)$
- Binary search:  $O(\log n)$
- DFS/BFS:  $O(\text{\#edges} + \text{\#vertices})$
- Hashtable with collision list:  $O(\text{max\_collision\_length})$
- Red-black tree search/insertion/deletion:  $O(\log n)$

# Space evaluation

---

A N-byte request at least consumes N-byte

Data structure overhead:

- Ex) Doubly-linked list: next and prev pointer (two words)

Internal fragmentation

- Ex) 3-Byte is requested, but 4-Byte is returned (1-Byte wasted)

External fragmentation

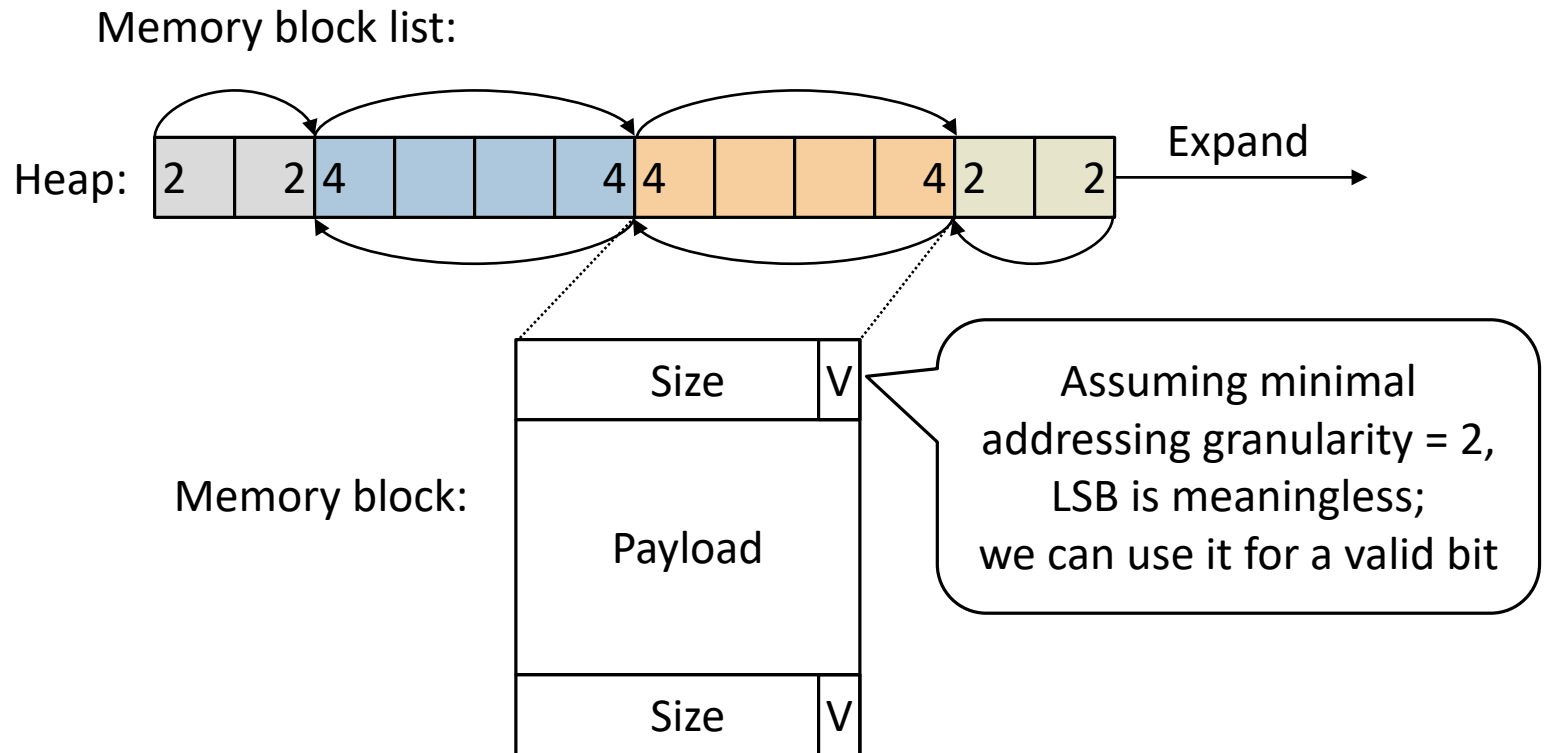
- Ex) There is total 4-Byte of free memory, but increased the heap to satisfy 4-Byte malloc request (4-Byte wasted)

# Example – Implicit free-list (IFL)

---

# Data structure


Summary: maintain a doubly linked list of memory blocks






# Operations


(1) Initial




(2) malloc(2); expand




(3) free(a); invalidate




(4-1) free(b); invalidate



(4-2) merge (a.k.a. coalesce)



(5) malloc(6); split



(You may find the first-fit or the best-fit)

# Implementation

---

malloc:

- Linearly search for an invalid memory block
- If nothing is found, expand the heap
- Mark the block as valid, and return the address

free:

- Mark the memory block as invalid
- Merge the adjacent blocks if they are also invalid
- ~~◦ If the last block becomes invalid, shrink the heap~~
  - For this assignment, we assume the heap never shrinks

# Performance

---

## Speed

- ☹️ malloc:  **$O(\text{\#memory\_blocks})$ ; linear search**
- ☺️ free:  $O(1)$ ; set invalid & coalesce

## Space

- ☺️ Overhead (2-word): next and prev displacement; valid bit may reuse next and prev's LSB
- ☺️ Internal fragmentation: fine
- External fragmentation:
  - ☹️ First-fit → **severe** / ☺️ Best-fit → fine

# Assignment

---

- Write a dynamic memory allocator for C programs
  - Functions including `mm_init`, `mm_malloc`, `mm_free`, `mm_realloc`
- Assure them to work correctly and efficiently
- Hand in only one source code file (`mm.c`) and your report
- For more details, please refer to the README